

Search that finds what an agent remembered


Keepsake ranks memory with BM25 over a posting table, in plain SQL under row-level security. These are the measurements behind that choice: public retrieval benchmarks, a 100k-concept load test, and Claude answering questions from its own memory.

0.682 SciFact nDCG@10, above published BM25 at 0.665 0.2.0: 0.387	41ms A 3-term query in a tenant of 100k concepts, timed as a superuser 0.2.0: 306 ms · app role: 63 ms	49/56 Held-out LongMemEval questions Claude answers from memory previous prompt: 44/56
--	---	---


RETRIEVAL QUALITY

Better ranking, in a fraction of the time

Each ranker answers the same queries over the same documents. Scores come from human relevance judgments. Up and to the left is better.



Every dataset, every metric



Keepsake BM25 (shipped)
 Alternatives we measured
 Published BM25, for reference

nDCG@10 scores the order of the top ten results. 1.0 puts every relevant document first.

METHOD

How to read and reproduce these numbers

- Retrieval runs on [BEIR](#) (SciFact, NFCorpus, FiQA) and on [LongMemEval S](#) session retrieval over 470 questions.
- Each agent question ran once. Judge noise alone flipped about 4 of 56 answers between runs, so a gap smaller than 5 is not significant.
- LongMemEval's histories come from public chat datasets. Real agent memory is shorter and more curated.

```
# each needs Docker and Go; agents a
python3 bench/beir.py
python3 bench/rankers.py
python3 bench/run.py
python3 bench/longmemeval.py
python3 bench/run.py --trials 1 \
  --tasks-file bench/data/longmemeval
```

Measured on 2026-09-23 and 2026-09-24 with the scripts in bench/ of [roee-fs/keepsake](#). BM25 parameters: k1 = 0.9, b = 0.4.